

The Valence Layer Algorithm: A Scalable Computational Method for Consecutive Prime Gaps and Integer Sequences

André Gualberto

July 15, 2026

Abstract

We present the Valence Layer Algorithm, a deterministic framework for generating consecutive primes and decomposing prime gaps using the integer sequence $\mathcal{V} = \{1000, 500, 200, \dots, 2\}$. The algorithm is implemented in Fortran (64/128-bit), C+GMP (arbitrary precision), and Python, scaling from desktop hardware to numbers with thousands of digits. Numerical analysis across five scales (10^3 to 10^{15}) yields a Hurst exponent $H \approx 0.5$, confirming that gaps behave as white noise with fractal dimension $D \approx 1.5$. We verify Goldbach's conjecture for even numbers up to 10^7 and compute the Liouville partial sums $L(N)$ up to 10^8 , observing $\max |L|/\sqrt{N} \approx 1.33$.

Key words and phrases: prime gap, valence layer algorithm, Liouville function, Goldbach conjecture, Hurst exponent.

2020 Mathematics Subject Classification: Primary 11A41; Secondary 11N05, 11P32, 11Y16.

1 Introduction

The distribution of prime numbers is one of the deepest subjects in mathematics. While the Prime Number Theorem (PNT) describes their asymptotic density $\pi(N) \sim N/\ln N$, the local structure—prime gaps—exhibits irregularity that is not yet fully understood.

The relations of the system are governed by a critical scaling factor of $1/2$. This exponent is the algebraic seed of the symmetry underlying the

algorithm: 2 is the fundamental building block of the valence toolbox, while $1/2$ represents the critical line $\Re(s) = 1/2$ of the non-trivial zeros of the Riemann Zeta function $\zeta(s)$. Under the Riemann Hypothesis, the error term in the prime distribution oscillates with an amplitude bound of $O(N^{1/2+\epsilon})$, a boundary mirrored by our empirical analysis of the Liouville partial sums $L(N) = O(N^{1/2+\epsilon})$ and the fractal fluctuations of the Hurst exponent ($H \approx 0.5$) describing the prime gaps.

This paper introduces the *Valence Layer Algorithm*, a deterministic method that:

1. Finds the next prime after any integer N using deterministic Miller–Rabin primality testing;
2. Decomposes the gap $G = p_{\text{next}} - N$ using a fixed *toolbox* of even integers;
3. Provides statistical analysis of gap distributions, Goldbach partitions, and the Liouville function.

2 The Valence Layer Algorithm

2.1 The Toolbox

Let $\mathcal{V} = \{v_1, v_2, \dots, v_k\}$ be a finite, ordered set of even integers sorted in descending order, defined as:

$$\mathcal{V} = \{1000, 500, 200, 120, 64, 34, 32, 30, 28, 24, 20, 18, 16, 12, 10, 8, 6, 4, 2\}.$$

Since $2 \in \mathcal{V}$, any even integer can be uniquely represented. We formalize the decomposition of a prime gap $G = p_{\text{next}} - N$ as an expansion over a state-space of valences. Let the valuation operator $\Phi_{\mathcal{V}} : \mathbb{E} \rightarrow \mathbb{Z}^k$ map any even integer G to a coefficient vector $\mathbf{s} = (s_1, s_2, \dots, s_k)$ such that:

$$G = \sum_{i=1}^k s_i \cdot v_i$$

where the coefficients $s_i \in \mathbb{N}$ are computed recursively under the greedy restriction constraint:

$$s_i = \left\lfloor \frac{G_{i-1}}{v_i} \right\rfloor, \quad G_i = G_{i-1} \pmod{v_i}$$

with the initial boundary condition $G_0 = G$.

Proposition 1 (Gap Parity). *Every gap between consecutive odd primes is even. The only exception is $3 - 2 = 1$.*

Proof. All primes $p > 2$ are odd. The difference of two odd numbers is even. For 2 and 3, the gap is 1 by inspection. $\square$

Corollary 2. *Every prime gap $G > 1$ can be decomposed using $\mathcal{V}$.*

Proposition 3. *The greedy expansion induced by $\Phi_{\mathcal{V}}$ terminates in finite steps and is exact for any prime gap $G > 1$.*

Proof. By Proposition 1, G is even. Since $v_k = 2$ is the infimum of $\mathcal{V}$, the remainder at the final step must satisfy $G_{k-1} \pmod{2} = 0$. Thus, G_{k-1} is exactly divisible by 2, yielding an exact decomposition with zero residual error. $\square$

2.2 Algorithm

Given N :

1. If $N < 2$, return 2.
2. If $N = 2$, return 3.
3. Start from the next odd candidate $c = N + 1$ (or $N + 2$ if N is odd).
4. Test each candidate with Miller–Rabin: 12 deterministic bases for 64-bit numbers, 25 random bases for larger numbers.
5. When a prime p is found, compute $G = p - N$ and decompose G using $\mathcal{V}$.

2.3 Implementation

The algorithm is implemented in three languages:

- **Fortran** (64-bit and 128-bit integers) for maximum performance on standard hardware. The 128-bit version finds the largest signed 128-bit prime $M_{127} = 2^{127} - 1$.
- **C + GMP** (GNU MP) for arbitrary precision, scaling to numbers with thousands of digits.
- **Python** with `gmpy2` for rapid prototyping and statistical analysis.

All source code is available at <https://github.com/angualberto/The-Valence-Layer-Algorithm>

3 Numerical Results

3.1 Carmichael Numbers

Carmichael numbers are composites that pass Fermat’s test but are correctly identified as composite by Miller–Rabin. Table 1 lists 16 Carmichael numbers tested.

Table 1: Carmichael numbers correctly rejected.

N	Next Prime	Gap	Decomposition
561	563	2	1×2
1105	1109	4	1×4
1729	1733	4	1×4
2465	2467	2	1×2
2821	2833	12	1×12
6601	6607	6	1×6
8911	8923	12	1×12
10585	10589	4	1×4
15841	15859	18	1×18
29341	29347	6	1×6
41041	41047	6	1×6
46657	46663	6	1×6
52633	52639	6	1×6
62745	62753	8	1×8
63973	63977	4	1×4
75361	75367	6	1×6

3.2 Benchmarks

Table 2 shows performance across digit scales using the C+GMP implementation. The average gap consistently follows $\ln N$ as predicted by the PNT.

The 5000-digit benchmark found the next prime in 1039 s (17.3 min).

3.3 Fractal Analysis

The gap distribution was analyzed across five scales (10^3 to 10^{15}). For each scale, 500 consecutive primes were generated using `gmpy2.next_prime`.

Table 2: Benchmark results (C + GMP).

Digits	Primes	Mean Gap	$\ln N$	Ratio
100	10,000	232.5	230.3	1.010
200	500	442.9	460.5	0.962
500	1	1300	1151	1.129
1000	1	13540	2303	5.88
2000	1	112	4605	0.024
3000	1	798	6908	0.116
5000	1	9978	11513	0.867

3.3.1 Hurst Exponent

The rescaled range (R/S) analysis gives the Hurst exponent H :

$$E[R(n)/S(n)] \sim C n^H, \quad n \rightarrow \infty.$$

$H = 0.5$ indicates white noise (uncorrelated increments), $H > 0.5$ persistence, and $H < 0.5$ anti-persistence.

Table 3: Hurst exponents across scales.

Scale	Mean Gap	$\ln N$	Ratio	H
10^3	8.0	6.9	1.16	0.456
10^6	13.3	13.8	0.96	0.501
10^9	20.6	20.7	0.99	0.574
10^{12}	28.0	27.6	1.01	0.550
10^{15}	39.2	34.5	1.14	0.599

The mean $H \approx 0.54$ is close to 0.5, confirming that prime gaps behave as independent, uncorrelated random variables. The fractal dimension $D = 2 - H \approx 1.46$.

3.3.2 Normalized Gap Distribution

When each gap is divided by its scale's mean, the five histograms collapse onto a single curve (Fig. 1), demonstrating auto-similarity.

3.4 Last Two Digits

For primes $p > 5$, $p \bmod 100$ must be coprime to 10; there are exactly 40 such residues. By Dirichlet's theorem, each should appear with frequency $1/40 = 2.5\%$.




Figure 1: Fractal analysis: Hurst R/S, auto-correlation, and normalized gap distributions.

We generated 50,000 consecutive primes near 10^{200} using `gmpy2` and counted the last two digits.

Table 4: Last-two-digit distribution for 50,000 primes near 10^{200} .

Statistic	Value
Mean count per residue	1250.0
Expected standard deviation	34.9
Observed standard deviation	30.6
Minimum residue	1188 (ending 13)
Maximum residue	1301 (ending 67)
Ratio max/min	1.095
χ^2 (39 df)	consistent with uniformity

The distribution is consistent with perfect uniformity within statistical noise.

4 Goldbach's Conjecture

Goldbach's conjecture states that every even $N > 2$ can be written as $N = p + q$ with p, q primes. Define the counting function $r(N) = \#\{(p, q) : p + q = N, p, q \text{ prime}\}$.

The Hardy–Littlewood asymptotic is:

$$r(N) \sim 2C_2 \frac{N}{\ln^2 N} \prod_{\substack{p|N \\ p>2}} \frac{p-1}{p-2}, \quad C_2 = \prod_{p>2} \left(1 - \frac{1}{(p-1)^2}\right) \approx 0.66016.$$

4.1 Numerical Verification

We tested even numbers up to 10^7 using `gmpy2.is_prime`.

Table 5: Goldbach partition counts.

N	$r(N)$	Asymptotic	Ratio
10^4	127	222.3	0.571
10^5	810	1372.4	0.590
10^6	5402	9372.0	0.576
10^7	38,807	68,300.7	0.568

Every even N tested has at least one Goldbach partition. The ratio $r(N)/r_{\text{asym}}(N)$ is stable at ≈ 0.57 , consistent with the asymptotic converging slowly from below.

5 The Liouville Function

The Liouville function is defined as $\lambda(n) = (-1)^{\Omega(n)}$, where $\Omega(n)$ is the total number of prime factors of n counted with multiplicity. The partial sums $L(N) = \sum_{n \leq N} \lambda(n)$ satisfy

$$\sum_{n=1}^{\infty} \frac{\lambda(n)}{n^s} = \frac{\zeta(2s)}{\zeta(s)}, \quad \Re(s) > 1.$$

We computed $L(N)$ up to $N = 10^8$ using a C sieve to count $\Omega(n)$. The implementation runs in 8 seconds on a standard desktop. Table 6 shows the behavior of $\max_{N \leq 10^8} |L(N)|/N^{1/2+\varepsilon}$.

Figure 2 shows $|L(N)|/N^{1/2+\varepsilon}$ for $\varepsilon = 0, 0.05, 0.10, 0.20$. For $\varepsilon = 0$ the ratio stabilises at ≈ 1.3 ; for any $\varepsilon > 0$ it decays to zero.

Figure 3 compares $|L(N)|$ with the envelopes $\sqrt{N}$ and $1.36\sqrt{N}$.

ε	$\max_{N \leq 10^8} L(N) /N^{1/2+\varepsilon}$	Behavior
0.00	1.231	Constant (≈ 1.3)
0.05	0.542	$\rightarrow 0$
0.10	0.249	$\rightarrow 0$
0.20	0.060	$\rightarrow 0$

Table 6: Ratio $\max |L(N)|/N^{1/2+\varepsilon}$ for different ε . For any $\varepsilon > 0$ the ratio tends to zero with N .




Figure 2: Ratio $|L(N)|/N^{1/2+\varepsilon}$ for $\varepsilon = 0, 0.05, 0.10, 0.20$.




Figure 3: $|L(N)|$ vs. $\sqrt{N}$ and $1.36\sqrt{N}$.

6 Bitwise Sieve Results in Fermat Factorization

6.1 Modular Residue Theory and the Fermat Bit Mask

The classical Fermat method seeks to solve the Diophantine equation $x^2 - y^2 = N$, which implies that the intermediate determinant $\Phi(x) = x^2 - N = y^2$ must be a perfect square over $\mathbb{Z}$.

We define the set of quadratic residues of a modulus m as:

$$\mathcal{Q}_m = \{a \in \mathbb{Z}_m \mid \exists k \in \mathbb{Z}_m \text{ such that } k^2 \equiv a \pmod{m}\}$$

For any choice of m , the necessary condition for a candidate x to produce a perfect square is that its projection satisfies:

$$\Phi(x) \pmod{m} \in \mathcal{Q}_m$$

For $m = 16$, the set of valid residues is extremely restricted: $\mathcal{Q}_{16} = \{0, 1, 4, 9\}$. We define the characteristic function of the filter $\chi_{\mathcal{Q}_{16}} : \mathbb{Z}_{16} \rightarrow \{0, 1\}$ associated with the control mask constant $\text{MASK}_{16} = 0\text{x}213_{16}$:

$$\chi_{\mathcal{Q}_{16}}(z) = (\lfloor \text{MASK}_{16} \cdot 2^{-z} \rfloor) \pmod{2}$$

At the processor register level, the transition evaluation reduces to the bitwise operation with cost $\mathcal{O}(1)$:

$$f(x) = (0\text{x}213 \gg ((x^2 - N) \& 15)) \& 1$$

If $f(x) = 0$, the candidate is rejected instantly without needing to invoke the floating-point operation $\sqrt{\cdot}$. For $m = 64$, the filter extends isomorphically using the set of residues $\mathcal{Q}_{64}$ of 12 elements mapped onto the 64-bit mask $\text{MASK}_{64} = 0\text{x}0282410120101113_{16}$.

For the odd coprime base sieves implemented in Algorithm 1, the rejection characteristic functions are derived analogously by mapping the quadratic residue sets:

$$\mathcal{Q}_9 = \{0, 1, 4, 7\} \subset \mathbb{Z}_9$$

$$\mathcal{Q}_5 = \{0, 1, 4\} \subset \mathbb{Z}_5$$

$$\mathcal{Q}_7 = \{0, 1, 2, 4\} \subset \mathbb{Z}_7$$

The control hexadecimal constants are generated by setting the bits corresponding to the allowed residues:

$$\begin{aligned}\text{MASK}_9 &= \sum_{r \in \mathcal{Q}_9} 2^r = 2^0 + 2^1 + 2^4 + 2^7 = 147 = 0\text{x}93_{16} \\ \text{MASK}_5 &= \sum_{r \in \mathcal{Q}_5} 2^r = 2^0 + 2^1 + 2^4 = 19 = 0\text{x}13_{16} \\ \text{MASK}_7 &= \sum_{r \in \mathcal{Q}_7} 2^r = 2^0 + 2^1 + 2^2 + 2^4 = 23 = 0\text{x}17_{16}\end{aligned}$$

Thus, the filtering operation in any base m is generalized in silicon by the logical relation with cost $\mathcal{O}(1)$:

$$f_m(y^2) = (\text{MASK}_m \gg (y^2 \pmod{m})) \& 1$$

Algorithm 1 Fermat Factorization with Stacked Bitwise Sieves

Require: N odd composite

Ensure: p, q prime factors of N

```

1:  $x \leftarrow \lceil \sqrt{N} \rceil$ 
2: if  $(N \& 3) = 1$  then
3:   if  $(x \& 1) = 0$  then
4:      $x \leftarrow x \mid 1$  ▷ forces odd parity via OR bitwise
5: else
6:   if  $(x \& 1) = 1$  then
7:      $x \leftarrow x + 1$  ▷ forces even parity
8: loop
9:    $y^2 \leftarrow x^2 - N$ 
10:  if  $(0x213 \gg (y^2 \& 15)) \& 1 = 0$  then
11:     $x \leftarrow x + 2$  ▷ rejected by Mod 16 sieve
12:    continue
13:  if  $(0x93 \gg (y^2 \pmod{9})) \& 1 = 0$  then
14:     $x \leftarrow x + 2$  ▷ rejected by Mod 9 sieve
15:    continue
16:  if  $(0x13 \gg (y^2 \pmod{5})) \& 1 = 0$  then
17:     $x \leftarrow x + 2$  ▷ rejected by Mod 5 sieve
18:    continue
19:   $r \leftarrow \lfloor \sqrt{y^2} \rfloor$  ▷ sqrt() restricted to survivors
20:  if  $r^2 = y^2$  then
21:    return  $p \leftarrow x - r, q \leftarrow x + r$ 
22:   $x \leftarrow x + 2$ 
```

Table 7 shows the impact of each quadratic residue sieve on the factorization of $N = 3999971 \times 4999999 = 19999851000029$ by the Fermat method. Each sieve eliminates candidates for x using only bit operations (&, shift, XOR) before calling $\sqrt{y^2}$.

Table 7: Candidate rejection by stacked bitwise sieves.

Sieves	Rejection	sqrt() calls	Reduction
No filter (only parity $N \& 3$)	0%	13 933	$1 \times$
Mod 16 (bitwise, $x \& 15$)	50.0%	6 967	$2 \times$
Mod 64 (bitwise, $x \& 63$)	75.0%	3 484	$4 \times$
Mod 16 + Mod 9	83.3%	2 323	$6 \times$
Mod 16 + Mod 9 + Mod 7	92.9%	996	$14 \times$
Mod 64 + Mod 9 + Mod 5	95.0%	698	$20 \times$
Mod 64 + Mod 9 + Mod 5 + Mod 7	97.9%	299	$47 \times$

6.2 Algebraic Modeling of the Carry Correlation Attack

The symmetric algorithm RC5 operates on the Cartesian product of truncated rings $\mathbb{Z}_{2^w} \times \mathbb{Z}_{2^w}$. We define the bitwise carry operator c_k for the addition of two machine words $X, Y \in \mathbb{Z}_{2^w}$ by the logical recurrence:

$$\begin{aligned} c_0 &= 0 \\ c_{k+1} &= (x_k \wedge y_k) \vee (x_k \wedge c_k) \vee (y_k \wedge c_k) \end{aligned}$$

When the rotation amount of a round falls into the identity state ($B \equiv 0 \pmod{w}$), the rotation endomorphism $\rho(A, B)$ degenerates, resulting in the linear additive simplification:

$$A' = (A \oplus B) + S[2i] \pmod{2^w}$$

Let $X = A \oplus B$ be a known and controllable input. For algebraic rigor, the differentiation operator $\frac{\partial A'}{\partial x_k}$ is formally defined as the discrete Boolean difference (Gibbs derivative) applied over the Galois field $\mathbb{Z}_{2^w}$:

$$\frac{\partial f(X)}{\partial x_k} = f(X) \oplus f(X \oplus 2^k)$$

This formulation guarantees that the variation of the carry bit directly maps the state transitions of the additive register. The discrete partial derivative of the output bit at position $k + 1$ with respect to the unit perturbation $\Delta X = 2^k$ is given by:

$$\frac{\partial A'}{\partial x_k} \implies s_k = c_{k+1} \oplus x_k \oplus (A')_k$$

This direct phase correlation allows the isolation of the carry term c_{k+1} and the linear inductive reconstruction of the subkey $S[2i]$ bits from right to left. The attack complexity is definitively reduced from an exponential barrier to a strictly linear sweep of word size:

$$\mathcal{O}(2^w) \longrightarrow \mathcal{O}(w)$$

Table 8 summarizes the practical gain for RC5-32/12/16.

Table 8: Carry propagation attack on RC5.

Method	Cost	Reduction
Brute force (2^{32} keys)	4 294 967 296 tests	$1\times$
Lock sieve (3.125%)	$\approx 1\,342\,177\,280$	$32\times$
Linear deduction by carry	32 additions	$134\,000\,000\times$

6.3 RC5 Structure and Rotation Lock

RC5- $w/r/b$ encrypts blocks of $2w$ bits using r rounds with data-dependent rotations. Each round applies:

Algorithm 2 RC5: single round encryption

Require: A, B (block halves), $S[2i], S[2i + 1]$ (subkeys)

Ensure: A', B' (next state)

- 1: $A \leftarrow ((A \oplus B) \lll B) + S[2i]$ $\triangleright$ rotation by $B \bmod w$
 - 2: $B \leftarrow ((B \oplus A) \lll A) + S[2i + 1]$
-

When $B \bmod w = 0$ (probability $1/w = 3,125\%$), the rotation **locks**: $\lll 0$ is identity. At that instant:

$$A' = (A \oplus B) + S[2i],$$

exposing the subkey $S[2i]$ to a differential attack.

6.4 RC5 Avalanche Test

Table 9 shows the avalanche effect of RC5-32/12/16: flipping 1 bit of the plaintext changes on average 32.58 of the 64 bits of the ciphertext (ideal: 32). The histogram (Fig. 4) has a Gaussian shape centered at 33 bits, with 10 of the 64 samples reaching exactly that value.

Table 9: Avalanche test for RC5-32/12/16 (64 samples).

Statistic	Value	Ideal
Mean bits changed	32.58	32.00
Standard deviation (population)	4.05	4.00
Minimum	21	—
Maximum	40	—
Absolute deviation from ideal	0.58	0.00

Bit change count (64 samples):

```

21 22      25 26      27 28 29 30 31 32 33 34 35 36 37 38 39 40
#  #          #  #    ### ##### ### #####  ### ##### ##### #####  #####  ### #

```

Figure 4: Histogram of the avalanche effect.

7 Gap Inversion Attack

7.1 Multiplicative Invariance of Gaps in Moduli

Consider the projection of two RSA moduli generated with shared factors $N_1 = p \cdot q_1$ and $N_2 = p \cdot q_2$, where q_1, q_2 are consecutive primes on the number line. The absolute difference between the moduli is given by the scaling relation:

$$\Delta N = |N_2 - N_1| = p \cdot |q_2 - q_1| = p \cdot \text{gap}_q$$

Since the distribution of consecutive prime gaps obeys the local scale $\text{gap}_q \ll \sqrt{N}$, the problem of factoring the composite number N_1 is directly mapped to the identification of the divisor p of ΔN through a search restricted to the valences of the gap subgroup $\mathcal{V}$:

$$p = \frac{\Delta N}{\sum_{i=1}^m K(s_i)}$$

This transforms the complexity barrier from sub-exponential to linear in the size of the observed gap $\mathcal{O}(g)$.

7.2 BigInt Implementation (Limb Concatenation)

For primes above 32 bits, $N = p \times q$ exceeds 2^{64} . We use a minimal big integer library by concatenation of 64-bit limbs (*little-endian*):

$$N = \sum_{i=0}^{n-1} L_i \cdot 2^{64i},$$

where $n = \lceil \text{bits}(N)/64 \rceil$. With $n \leq 64$ we cover up to 4096 bits (RSA-2048). The required operations are:

1. Subtraction ($\Delta N = N_2 - N_1$);
2. Modulo by g (64 bits): $\Delta N \bmod g$;
3. Division by g : $p = \Delta N/g$;
4. BigInt modulo: $N_1 \bmod p$ (verification).

7.3 Experimental Results

Table 10 shows the attack performance for primes from 8 to 256 bits. The number of steps depends only on the gap g , not on the key size.

Table 10: Gap inversion attack: steps vs. prime size.

Bits(p)	p (hex)	g	Steps	Status
8	0x83	2	1	OK
16	0x8003	4	2	OK
24	0x800009	4	2	OK
32	0x8000000b	20	10	OK
40	0x8000000017	6	3	OK
48	0x800000000005	32	16	OK
56	0x80000000000003	40	20	OK
256	0xdbff...373f	146	73	OK

7.4 Comparison with Other Methods

Table 11 compares the gap attack with trial division and Fermat for $N = 19999851000029$ (32-bit primes). The gap attack is the only one whose complexity does not depend on the prime size.

7.5 Limitation

The attack requires two moduli that share the same prime p . For a single isolated N , trial division (+2) is the fastest method for keys up to 32 bits

Table 11: Comparison of factorization methods.

Method	Iterations	Time (C)	Dependency
Trial division (+2)	1 999 985	0.023 s	$O(\sqrt{N})$
Fermat + parity	13 933	0.005 s	$O(p - q)$
Sieve + OMP	272 136	0.060 s	$O(\sqrt{N} \log \log \sqrt{N})$
Gap inversion (2 keys)	1–250	$< 1 \mu\text{s}$	$O(g)$, $g < 500$
Classic GCD (2 keys)	1	$< 1 \mu\text{s}$	requires identical p

(0.023 s). For real RSA keys (≥ 1024 bits), no trial division method is viable—RSA security remains intact for well-generated keys.

The theoretical contribution of the gap attack is to demonstrate that the structure of gaps between primes is preserved multiplicatively in the moduli, and that $\Delta N = p \times g$ reduces the factorization problem to a search over small gaps ($O(\log g)$ instead of $O(\sqrt{N})$).

8 Conclusion

The Valence Layer Algorithm provides:

1. A practical, scalable method for finding consecutive primes and decomposing their gaps into the integer sequence $\{1000, 500, 200, \dots, 2\}$;
2. Numerical confirmation of the PNT, auto-similarity of gaps ($H \approx 0.5$), and uniform last-digit distribution across 40 residue classes;
3. Verification of Goldbach’s conjecture up to 10^7 and a structural connection through parity and the Hardy–Littlewood circle method;
4. Computation of the Liouville function $L(N)$ up to 10^8 , with $\max |L(N)|/\sqrt{N} \approx 1.33$ stable across three scales.

All code and data are publicly available.

References

- [1] E. C. Titchmarsh and D. R. Heath-Brown, *The Theory of the Riemann Zeta Function*, 2nd ed., Oxford University Press, 1986.
- [2] G. H. Hardy and J. E. Littlewood, Some problems of ‘Partitio Numerorum’; III: On the expression of a number as a sum of primes, *Acta Math.* **44** (1923), 1–70.

- [3] B. Efron and T. Hastie, *Computer Age Statistical Inference*, Cambridge University Press, 2016.
- [4] M. O. Rabin, Probabilistic algorithm for testing primality, *J. Number Theory* **12** (1980), 128–138.